

Machine Learning

Dmytro Fishman (dmytro.fishman@ut.ee)



UNIVERSITY OF TARTU
Institute of Computer Science


BIIT


September 13, 2021


Deadlines

Assignment	Date of assignment	Deadline (midnight 23:59)
HW1	Sep 6	Sep 19
HW2	Sep 20	Oct 3
HW3	Oct 4	Oct 17
Paper summary	Oct 11	Oct 24
HW4	Oct 18	Oct 31
HW5	Nov 1	Nov 14
HW6	Nov 22	Dec 5
Project	TBA	Dec 13 - 15


***All deadlines are subject to change, check out CampusWire and website for updates**







Classification



Regression

Machine
Learning

Unsupervised
Learning

Reinforcement
Learning

Classification


Supervised
Learning

Regression


Machine
Learning

Unsupervised
Learning


Reinforcement
Learning




Binary Classification




Multiclass Classification



Nearest Neighbour Classifier



K-Nearest Neighbour Classifier



0 $\longleftrightarrow$? $\longleftrightarrow$ 1

Classification


Supervised
Learning

Regression


Machine
Learning

Unsupervised
Learning


Reinforcement
Learning



Linear Regression




Linear Regression



Linear Regression (example)

$$\hat{y} = 2.2 + .6x$$



x	y	x - $\bar{x}$	y - $\bar{y}$	(x - $\bar{x}$) ²	(x - $\bar{x}$)(y - $\bar{y}$)
1	2	-2	-2	4	4
2	4	-1	0	1	0
3	5	0	1	0	0
4	4	1	0	1	0
5	5	2	1	4	2

$$\bar{x} = 3 \quad \bar{y} = 4$$

$$10$$

$$6$$

$$w_1 = \frac{\sum (x - \bar{x})(y - \bar{y})}{\sum (x - \bar{x})^2} = \frac{6}{10} = .6$$


$$4 = w_0 + .6 * 3$$

$$w_0 = 2.2$$

Linear Regression (example)

$$\hat{y} = w_0 + w_1 x$$

Check out the video for how the formulas are derived!



Derivation of OLS estimates of slope and intercept

slope: $\hat{\beta}_1 = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2} = \frac{s_{xy}}{s_x^2}$

intercept: $\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}$

Created with Doceri

$$w_1 = \frac{\sum (x - \bar{x})(y - \bar{y})}{\sum (x - \bar{x})^2}$$

$$w_0 = \bar{y} - w_1 \bar{x}$$


$$\bar{x} = 3 \quad \bar{y} = 4$$

$$10$$


$$6$$

<https://youtu.be/jqoHefilf9U>


Linear Regression




Linear Regression



Linear Regression





Linear Regression



Decision Tree Algorithm



Decision Tree Algorithm

By asking a simple **question** about value of **independent variable** it tries to predict a value of **dependent variable**




Decision Tree Algorithm

By asking a simple **question** about value of **independent variable** it tries to predict a value of **dependent variable**




Decision Tree Algorithm

By asking a simple **question** about value of **independent variable** it tries to predict a value of **dependent variable**





Decision Tree Algorithm

By asking a simple **question** about value of **independent variable** it tries to predict a value of **dependent variable**



Decision Tree Algorithm



By asking a simple **question** about value of **independent variable** it tries to predict a value of **dependent variable**



Decision Tree Algorithm

Here, **X** may correspond to any **vertical** line.



For example if $X = 2.5$:



Decision Tree Algorithm

What would be MSE if **Y = 4** and **Z = 5**?



For example if $X = 2.5$:



What are most reasonable values for **Y** and **Z** (that minimise total MSE)?

Decision Tree Algorithm



$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$



What are most reasonable values for **Y** and **Z** (that minimise total MSE)?

Decision Tree Algorithm



$$MSE = \frac{1}{n} \sum_{i=1}^n (\overset{\text{real value}}{y_i} - \underset{\text{predicted value}}{\hat{y}_i})^2$$



What are most reasonable values for **Y** and **Z** (that minimise total MSE)?

Decision Tree Algorithm



$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \frac{(y_1 - \hat{y}_1)^2 + (y_2 - \hat{y}_2)^2 + (y_3 - \hat{y}_3)^2 + (y_4 - \hat{y}_4)^2 + (y_5 - \hat{y}_5)^2}{5}$$



What are most reasonable values for **Y** and **Z** (that minimise total MSE)?

Decision Tree Algorithm



$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \frac{(2 - 4)^2 + (y_2 - \hat{y}_2)^2 + (y_3 - \hat{y}_3)^2 + (y_4 - \hat{y}_4)^2 + (y_5 - \hat{y}_5)^2}{5}$$



What are most reasonable values for **Y** and **Z** (that minimise total MSE)?

Decision Tree Algorithm



$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \frac{(2 - 4)^2 + (y_2 - \hat{y}_2)^2 + (y_3 - \hat{y}_3)^2 + (y_4 - \hat{y}_4)^2 + (y_5 - \hat{y}_5)^2}{5}$$



What are most reasonable values for **Y** and **Z** (that minimise total MSE)?

Decision Tree Algorithm



$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \frac{(2 - 4)^2 + (4 - 4)^2 + (y_3 - \hat{y}_3)^2 + (y_4 - \hat{y}_4)^2 + (y_5 - \hat{y}_5)^2}{5}$$



What are most reasonable values for **Y** and **Z** (that minimise total MSE)?

Decision Tree Algorithm



$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \frac{(2 - 4)^2 + (4 - 4)^2 + (y_3 - \hat{y}_3)^2 + (y_4 - \hat{y}_4)^2 + (y_5 - \hat{y}_5)^2}{5}$$



What are most reasonable values for **Y** and **Z** (that minimise total MSE)?

Decision Tree Algorithm



$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \frac{(2 - 4)^2 + (4 - 4)^2 + (5 - 5)^2 + (y_4 - \hat{y}_4)^2 + (y_5 - \hat{y}_5)^2}{5}$$



What are most reasonable values for **Y** and **Z** (that minimise total MSE)?

Decision Tree Algorithm



$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \frac{(2 - 4)^2 + (4 - 4)^2 + (5 - 5)^2 + (y_4 - \hat{y}_4)^2 + (y_5 - \hat{y}_5)^2}{5}$$



What are most reasonable values for **Y** and **Z** (that minimise total MSE)?

Decision Tree Algorithm



$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \frac{(2 - 4)^2 + (4 - 4)^2 + (5 - 5)^2 + (4 - 5)^2 + (y_5 - \hat{y}_5)^2}{5}$$



What are most reasonable values for **Y** and **Z** (that minimise total MSE)?

Decision Tree Algorithm



$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \frac{(2 - 4)^2 + (4 - 4)^2 + (5 - 5)^2 + (4 - 5)^2 + (5 - 5)^2}{5}$$



What are most reasonable values for **Y** and **Z** (that minimise total MSE)?

Decision Tree Algorithm



$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \frac{(2 - 4)^2 + (4 - 4)^2 + (5 - 5)^2 + (4 - 5)^2 + (5 - 5)^2}{5}$$



What are most reasonable values for **Y** and **Z** (that minimise total MSE)?

Decision Tree Algorithm



$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \frac{(-2)^2 + (0)^2 + (0)^2 + (-1)^2 + (0)^2}{5}$$



What are most reasonable values for **Y** and **Z** (that minimise total MSE)?

Decision Tree Algorithm



$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \frac{4 + 0 + 0 + 1 + 0}{5}$$



What are most reasonable values for **Y** and **Z** (that minimise total MSE)?

Decision Tree Algorithm



$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \frac{5}{5}$$



What are most reasonable values for **Y** and **Z** (that minimise total MSE)?

Decision Tree Algorithm



$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = 1$$



What are most reasonable values for **Y** and **Z** (that minimise total MSE)?

Decision Tree Algorithm

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = 1 \quad \text{so, if } \mathbf{X} = \mathbf{2.5}, \mathbf{Y} = \mathbf{4} \text{ and } \mathbf{Z} = \mathbf{5}, \text{ MSE is } \mathbf{1}$$





What are most reasonable values for **Y** and **Z** (that minimise total MSE)?

Decision Tree Algorithm

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = 1 \quad \text{so, if } \mathbf{X} = \mathbf{2.5}, \mathbf{Y} = \mathbf{4} \text{ and } \mathbf{Z} = \mathbf{5}, \text{ MSE is } \mathbf{1}$$



Can we find better **Y** and **Z**?



What are most reasonable values for **Y** and **Z** (that minimise total MSE)?

Decision Tree Algorithm



$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \frac{(y_1 - \hat{y}_1)^2 + (y_2 - \hat{y}_2)^2 + (y_3 - \hat{y}_3)^2 + (y_4 - \hat{y}_4)^2 + (y_5 - \hat{y}_5)^2}{5}$$



What are most reasonable values for **Y** and **Z** (that minimise total MSE)?

Decision Tree Algorithm



$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \frac{(2 - 3)^2 + (4 - 3)^2 + (5 - 5)^2 + (4 - 5)^2 + (5 - 5)^2}{5}$$



What are most reasonable values for **Y** and **Z** (that minimise total MSE)?

Decision Tree Algorithm



$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \frac{1 + 1 + 0 + 1 + 0}{5}$$



What are most reasonable values for **Y** and **Z** (that minimise total MSE)?

Decision Tree Algorithm

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \frac{3}{5} = 0.6$$





What are most reasonable values for **Y** and **Z** (that minimise total MSE)?

Decision Tree Algorithm

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \frac{3}{5} = 0.6 \quad \text{so, if } \mathbf{X} = \mathbf{2.5}, \mathbf{Y} = \mathbf{3} \text{ and } \mathbf{Z} = \mathbf{5},$$



MSE is **0.6**



What are most reasonable values for **Y** and **Z** (that minimise total MSE)?

Decision Tree Algorithm



$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \frac{(2 - 3)^2 + (4 - 3)^2 + (5 - 4.66)^2 + (4 - 4.66)^2 + (5 - 4.66)^2}{5}$$



What are most reasonable values for **Y** and **Z** (that minimise total MSE)?

Decision Tree Algorithm

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \frac{1 + 1 + 0.12 + 0.43 + 0.12}{5}$$





What are most reasonable values for **Y** and **Z** (that minimise total MSE)?

Decision Tree Algorithm

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \frac{2.67}{5} = 0.53$$



so, if **Y = 3** and **Z = 4.66**,
MSE is **smallest**



Decision Tree Algorithm



$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \frac{2.67}{5} = 0.53$$

so, if **Y = 3** and **Z = 4.66**,
MSE is **smallest**





Decision Tree Algorithm

Hold on, how did we choose this **split** on the first place?





Decision Tree Algorithm

Hold on, how did we choose this **split** on the first place?
Maybe there are better options?





Decision Tree Algorithm

What are the possible **split** options in this case?





Decision Tree Algorithm

What are the possible **split** options in this case?





Decision Tree Algorithm

What are the possible **split** options in this case?





Decision Tree Algorithm

What are the possible **split** options in this case?





Decision Tree Algorithm

Are these meaningful?





Decision Tree Algorithm

How to compare remaining?





Decision Tree Algorithm

How to compare remaining?
For each one we can compute MSE





Decision Tree Algorithm



How to compare remaining?
For each one we can compute MSE





Decision Tree Algorithm



Decision Tree Algorithm





Decision Tree Algorithm





Decision Tree Algorithm



How to compare remaining?
For each one we can compute MSE



Decision Tree Algorithm





Decision Tree Algorithm





Decision Tree Algorithm



How to compare remaining?
For each one we can compute MSE



Decision Tree Algorithm





Decision Tree Algorithm





Decision Tree Algorithm

How to compare remaining?
For each one we can compute MSE





Decision Tree Algorithm

We choose the **split** that minimises total MSE





Decision Tree Algorithm

Thus, the resulting tree:





Decision Tree Algorithm

Can we make our decision tree more accurate?




Decision Tree Algorithm

Can we make our decision tree more accurate?
Yes, by going **deeper**!




Decision Tree Algorithm

What is the best value for **Z** in this case?





Decision Tree Algorithm

What is the best value for **Z** in this case?





Decision Tree Algorithm





Decision Tree Algorithm

What is the best value for **Y** in this case?





Decision Tree Algorithm

What is the best value for **Y** in this case?





Decision Tree Algorithm

What is the best value for **Y** in this case?





Decision Tree Algorithm

What is the best value for **Y** in this case?





Decision Tree Algorithm





Decision Tree Algorithm

In principle nobody can stop us from...




Decision Tree Algorithm

In principle nobody can stop us from...




Decision Tree Algorithm

In principle nobody can stop us from...




Every observation is in its own leaf of the tree

MSE = 0.0






Decision Tree Algorithm



You can think about this as super **powerful** model that predicts all training example **perfectly**




Every observation is in its own leaf of the tree








Complex, but ideal



Simple, but imperfect







VS



Complex, but ideal



Simple, but imperfect




How do we know which one to choose?




VS



Complex, but ideal




Simple, but imperfect




Train/val split


Initial dataset




Train/val split



Train/val split



Train/val split



Train/val split



Simple, but
imperfect

Complicated,
but ideal


Randomly
select 60%

Randomly
select 40%

Initial dataset



Train dataset



Validation (val) dataset

Train/val split



Simple, but
imperfect

Complicated,
but ideal

Randomly
select 60%


Randomly
select 40%

Initial dataset



MSE = 0.0

Validation (val) dataset



Train/val split




Simple, but
imperfect

Complicated,
but ideal

Randomly
select 60%

Randomly
select 40%

Initial dataset




Train/val split

Simple, but
imperfect



Complicated,
but ideal

Initial dataset



Randomly
select 60%

Randomly
select 40%



Train/val split




Simple, but
imperfect

Complicated,
but ideal

Randomly
select 60%

Randomly
select 40%

Initial dataset



Train/val split




Simple, but
imperfect

Complicated,
but ideal

Randomly
select 60%

Randomly
select 40%

Initial dataset



Train/val split


Simple, but
imperfect

Complicated,
but ideal


Randomly
select 60%

Randomly
select 40%


Initial dataset



Train dataset



Validation (val) dataset




Train/val split

Simple, but
imperfect



Complicated,
but ideal

Initial dataset



Randomly
select 60%

Randomly
select 40%



Train/val split

Simple, but
imperfect

Complicated,
but ideal


The **gap** between **training** and
validation performance is not large

Random
select 60%


Randomly
select 40%

MSE = 1.0

MSE = 0.5



MSE = 0.0



MSE = 2.5

Train/val split

Simple, but
imperfect

Complicated,
but ideal



Large gap between
training and validation

Randomly
select 60%

Randomly
select 40%

MSE = 0.0

MSE = 2.5



Overfitting


Simple, but
imperfect

Complicated,
but ideal

Large gap between
training and validation


Randomly
select 60%

MSE = 0.0



Randomly
select 40%

MSE = 2.5




Train/val split

Simple, but
imperfect



Complicated,
but ideal

Initial dataset




Randomly
select 60%

Randomly
select 40%



Train/val split



What if the split was **lucky** and validation set turned out to be too simple for one of the models?

Cross Validation (CV) Algorithm

Cross Validation (CV) Algorithm

The whole dataset **100%**

Cross Validation (CV) Algorithm

Training data **80%**


Test **20%**

Cross Validation (CV) Algorithm

Training data **80%**

Test **20%**


Test set is used later for
the **final evaluation**




Cross Validation (CV) Algorithm

Training data **80%**

Cross Validation (CV) Algorithm



Cross Validation (CV) Algorithm



Training data **80%**



Train on **60%** of data Validate on
20%

Cross Validation (CV) Algorithm




Training data **80%**




Train on **60%** of data Validate on
20%

Cross Validation (CV) Algorithm




Cross Validation (CV) Algorithm



Cross Validation (CV) Algorithm




Training data **80%**



Cross Validation (CV) Algorithm




Training data **80%**



Cross Validation (CV) Algorithm


Simple, but
imperfect



Average MSE = **0.65**


Cross Validation (CV) Algorithm

Simple, but
imperfect



Average MSE = **0.65**


Complicated,
but ideal



Average MSE = 1.4


Cross Validation (CV) Algorithm

Simple, but
imperfect



Average MSE = **0.65**

Complicated,
but ideal




Average MSE = 1.4

Choose the **best** model

Cross Validation (CV) Algorithm

Simple, but
imperfect

Complicated,
but ideal




Cross Validation (CV) Algorithm

Simple, but
imperfect

Choose a **test set**
(20% at random)


Complicated,
but ideal



Cross Validation (CV) Algorithm

Simple, but
imperfect

Complicated,
but ideal




Cross Validation (CV) Algorithm

Simple, but
imperfect

We have **4 folds** by
20% of data each

Complicated,
but ideal




Cross Validation (CV) Algorithm

Simple, but
imperfect

We have **4 folds** by
20% of data each

Complicated,
but ideal




Every iteration one of this points
will be a **validation set**

Cross Validation (CV) Algorithm

Simple, but
imperfect

Complicated,
but ideal

Iteration I




Cross Validation (CV) Algorithm

Simple, but
imperfect

Complicated,
but ideal

Iteration I




Cross Validation (CV) Algorithm

Simple, but
imperfect

Complicated,
but ideal

Iteration I

	Training	Val
I		
II		
III		
IV		
Average		




Cross Validation (CV) Algorithm

Simple, but
imperfect

	Training	Val
I		
II		
III		
IV		
Average		

Iteration I

Complicated,
but ideal




Cross Validation (CV) Algorithm

Simple, but
imperfect

Complicated,
but ideal

Iteration I

	Training	Val
I		
II		
III		
IV		
Average		




Cross Validation (CV) Algorithm

Simple, but
imperfect

Complicated,
but ideal

Iteration I

	Training	Val
I		
II		
III		
IV		
Average		




Cross Validation (CV) Algorithm

Simple, but
imperfect

Complicated,
but ideal

Iteration I

	Training	Val
I	0.66	
II		
III		
IV		
Average		




Cross Validation (CV) Algorithm

Simple, but
imperfect

Complicated,
but ideal

Iteration I

	Training	Val
I	0.66	
II		
III		
IV		
Average		




Cross Validation (CV) Algorithm

Simple, but
imperfect

Complicated,
but ideal

Iteration I

	Training	Val
I	0.66	
II		
III		
IV		
Average		




Cross Validation (CV) Algorithm

Simple, but
imperfect

Complicated,
but ideal

Iteration I

	Training	Val
I	0.66	1
II		
III		
IV		
Average		




Cross Validation (CV) Algorithm

Simple, but
imperfect

	Training	Val
I	0.66	1
II		
III		
IV		
Average		

Iteration I

Complicated,
but ideal




Cross Validation (CV) Algorithm

Simple, but
imperfect

	Training	Val
I	0.66	1
II		
III		
IV		
Average		

Iteration I



Complicated,
but ideal


	Training	Val
I		
II		
III		
IV		
Average		

Cross Validation (CV) Algorithm

Simple, but
imperfect

	Training	Val
I	0.66	1
II		
III		
IV		
Average		

Iteration I



Complicated,
but ideal


	Training	Val
I		
II		
III		
IV		
Average		

Cross Validation (CV) Algorithm

Simple, but
imperfect

	Training	Val
I	0.66	1
II		
III		
IV		
Average		

Iteration I



Complicated,
but ideal


	Training	Val
I		
II		
III		
IV		
Average		

Cross Validation (CV) Algorithm

Simple, but
imperfect

	Training	Val
I	0.66	1
II		
III		
IV		
Average		

Iteration I



Complicated,
but ideal


	Training	Val
I	0	
II		
III		
IV		
Average		

Cross Validation (CV) Algorithm

Simple, but
imperfect

	Training	Val
I	0.66	1
II		
III		
IV		
Average		

Iteration I



Complicated,
but ideal


	Training	Val
I	0	
II		
III		
IV		
Average		

Cross Validation (CV) Algorithm

Simple, but
imperfect

	Training	Val
I	0.66	1
II		
III		
IV		
Average		

Iteration I



Complicated,
but ideal

$(2^2)/1$


	Training	Val
I	0	
II		
III		
IV		
Average		

Cross Validation (CV) Algorithm

Simple, but
imperfect

	Training	Val
I	0.66	1
II		
III		
IV		
Average		

Iteration I



Complicated,
but ideal

$(2^2)/1$


	Training	Val
I	0	4
II		
III		
IV		
Average		

Cross Validation (CV) Algorithm

Simple, but
imperfect

	Training	Val
I	0.66	1
II		
III		
IV		
Average		

Iteration I



Complicated,
but ideal


	Training	Val
I	0	4
II		
III		
IV		
Average		

Cross Validation (CV) Algorithm

Simple, but
imperfect

	Training	Val
I	0.66	1
II		
III		
IV		
Average		

Iteration II



Complicated,
but ideal


	Training	Val
I	0	4
II		
III		
IV		
Average		

Cross Validation (CV) Algorithm

Simple, but
imperfect

	Training	Val
I	0.66	1
II		
III		
IV		
Average		

Iteration II



Complicated,
but ideal


	Training	Val
I	0	4
II		
III		
IV		
Average		

Cross Validation (CV) Algorithm

Simple, but
imperfect

	Training	Val
I	0.66	1
II	0.37	0.25
III		
IV		
Average		

Iteration II



Complicated,
but ideal


	Training	Val
I	0	4
II		
III		
IV		
Average		

Cross Validation (CV) Algorithm

Simple, but
imperfect

	Training	Val
I	0.66	1
II	0.37	0.25
III		
IV		
Average		

Iteration II



Complicated,
but ideal


	Training	Val
I	0	4
II		
III		
IV		
Average		

Cross Validation (CV) Algorithm

Simple, but
imperfect

	Training	Val
I	0.66	1
II	0.37	0.25
III		
IV		
Average		

Iteration II



Complicated,
but ideal


	Training	Val
I	0	4
II	0	1.25
III		
IV		
Average		

Cross Validation (CV) Algorithm

Simple, but
imperfect

	Training	Val
I	0.66	1
II	0.37	0.25
III		
IV		
Average		

Iteration III



Complicated,
but ideal


	Training	Val
I	0	4
II	0	1.25
III		
IV		
Average		

Cross Validation (CV) Algorithm

Simple, but
imperfect

	Training	Val
I	0.66	1
II	0.37	0.25
III		
IV		
Average		

Iteration III



Complicated,
but ideal


	Training	Val
I	0	4
II	0	1.25
III		
IV		
Average		

Cross Validation (CV) Algorithm

Simple, but
imperfect

	Training	Val
I	0.66	1
II	0.37	0.25
III	0.014	1.25
IV		
Average		

Iteration III



Complicated,
but ideal


	Training	Val
I	0	4
II	0	1.25
III		
IV		
Average		

Cross Validation (CV) Algorithm

Simple, but
imperfect

	Training	Val
I	0.66	1
II	0.37	0.25
III	0.014	1.25
IV		
Average		

Iteration III




Complicated,
but ideal

	Training	Val
I	0	4
II	0	1.25
III		
IV		
Average		

Cross Validation (CV) Algorithm

Simple, but imperfect

	Training	Val
I	0.66	1
II	0.37	0.25
III	0.014	1.25
IV		
Average		

Iteration **III**

Complicated,
but ideal


	Training	Val
I	0	4
II	0	1.25
III	0	1.25
IV		
Average		

Cross Validation (CV) Algorithm

Simple, but
imperfect

	Training	Val
I	0.66	1
II	0.37	0.25
III	0.014	1.25
IV		
Average		

Iteration **IV**



Complicated,
but ideal


	Training	Val
I	0	4
II	0	1.25
III	0	1.25
IV		
Average		

Cross Validation (CV) Algorithm

Simple, but
imperfect

	Training	Val
I	0.66	1
II	0.37	0.25
III	0.014	1.25
IV	1	0
Average		

Iteration **IV**



Complicated,
but ideal


	Training	Val
I	0	4
II	0	1.25
III	0	1.25
IV		
Average		

Cross Validation (CV) Algorithm

Simple, but
imperfect

	Training	Val
I	0.66	1
II	0.37	0.25
III	0.014	1.25
IV	1	0
Average		

Iteration **IV**




Complicated,
but ideal

	Training	Val
I	0	4
II	0	1.25
III	0	1.25
IV	0	4
Average		

Cross Validation (CV) Algorithm

Simple, but
imperfect

	Training	Val
I	0.66	1
II	0.37	0.25
III	0.014	1.25
IV	1	0
Average		




Complicated,
but ideal

	Training	Val
I	0	4
II	0	1.25
III	0	1.25
IV	0	4
Average		

Cross Validation (CV) Algorithm

Simple, but
imperfect

	Training	Val
I	0.66	1
II	0.37	0.25
III	0.014	1.25
IV	1	0
Average	0.511	0.625




Complicated,
but ideal

	Training	Val
I	0	4
II	0	1.25
III	0	1.25
IV	0	4
Average	0	2.65

Cross Validation (CV) Algorithm

Simple, but
imperfect

	Training	Val
I	0.66	1
II	0.37	0.25
III	0.014	1.25
IV	1	0
Average	0.511	0.625



Complicated,
but ideal


	Training	Val
I	0	4
II	0	1.25
III	0	1.25
IV	0	4
Average	0	2.65

Choose **the best** model based on **average validation**
performance across CV folds

Cross Validation (CV) Algorithm

Simple, but
imperfect

	Training	Val
I	0.66	1
II	0.37	0.25
III	0.014	1.25
IV	1	0
Average	0.511	0.625




Choose **the best** model based on **average validation**
performance across CV folds

Cross Validation (CV) Algorithm

Simple, but
imperfect

The **final model** then
should be trained on all
training examples

	Training	Val
I	0.66	1
II	0.37	0.25
III	0.014	1.25
IV	1	0
Average	0.511	0.625




Choose **the best** model based on **average validation**
performance across CV folds

Cross Validation (CV) Algorithm

Simple, but
imperfect

	Training	Val
I	0.66	1
II	0.37	0.25
III	0.014	1.25
IV	1	0
Average	0.511	0.625


The **final model** then
should be trained on all
training examples




The **test set** is used
for the final evaluation

Choose **the best** model based on **average validation**
performance across CV folds

Disclaimer: never evaluate models on 1 data point.
Here we show it only for the illustration purposes.



Disclaimer: never evaluate models on 1 data point.
Here we show it only for the illustration purposes.



Disclaimer 2: simple models are not always better than complex. Just make sure you don't overfit your data


Cross Validation (CV) Algorithm

For each parameter (e.g. **K** in K-NN) choice - new loop of CV

Cross Validation (CV) Algorithm


For each parameter (e.g. **K** in K-NN) choice - new loop of CV

$K = 1$




Average acc. = **0.87**

$K = 5$



Average acc. = 0.81

$K = 10$




Average acc. = 0.83

Cross Validation (CV) Algorithm


For each parameter (e.g. **K** in K-NN) choice - new loop of CV

$K = 1$




Average acc. = **0.87**

$K = 5$



Average acc. = 0.81

$K = 10$




Average acc. = 0.83

Before each CV loop we **shuffle** the data
to get new **random** folds




So far...

K-Nearest Neighbour Classifier





Linear Regression



Decision Tree Algorithm

We choose the **split** that minimises total MSE




Train/val split

Simple, but
imperfect



Complicated,
but ideal

Initial dataset




Randomly
select 60%

Randomly
select 40%




Cross Validation (CV) Algorithm

Simple, but
imperfect



Average MSE = **0.65**

Complicated,
but ideal



Average MSE = 1.4

Before each CV loop we **shuffle** the data
to get new **random** folds




Supervised Learning pipeline


Supervised Learning pipeline

1 Acquire Data


Supervised Learning pipeline




Supervised Learning pipeline




Supervised Learning pipeline




Supervised Learning pipeline




Supervised Learning pipeline



Supervised Learning pipeline



Supervised Learning pipeline





If I have seen further than others, it is by standing
upon the shoulders of giants.

(Isaac Newton)

References



- Machine Learning by **Andrew Ng** (<https://www.coursera.org/learn/machine-learning>)
- Introduction to Machine Learning by **Pascal Vincent** given at Deep Learning Summer School, Montreal 2015 (http://videolectures.net/deeplearning2015_vincent_machine_learning/)
- Welcome to Machine Learning by **Konstantin Tretyakov** delivered at AACIMP Summer School 2015 (<http://kt.era.ee/lectures/aacimp2015/1-intro.pdf>)
- Stanford CS class: Convolutional Neural Networks for Visual Recognition by **Andrej Karpathy** (<http://cs231n.github.io/>)
- Data Mining Course by **Jaak Vilo** at University of Tartu (https://courses.cs.ut.ee/MTAT.03.183/2017_spring/uploads/Main/DM_05_Clustering.pdf)
- Machine Learning Essential Concepts by **Ilya Kuzovkin** (<https://www.slideshare.net/iljakuzovkin>)
- From the brain to deep learning and back by **Raul Vicente Zafra** and **Ilya Kuzovkin** (<http://www.uttv.ee/naita?id=23585&keel=eng>)

That's all Folks!

BIIT



www.slideshare.net/DimaFishman



www.slideshare.net/DimaFishman

